

Deep Learning CS7015 Programming Assignment 1 Backpropagation

Shubhangi Ghosh - EE15B129
Monisha J - CS15B053

February 2018

1 Model

The implementation of neural network was done using python and numpy.

The important functions include :

forward_propagation()

back_propagation()

gradient_descent()

2 Supported Hyperparameters

The activation functions supported are *sigmoid*, *tanh* and *relu*.

The final layer activation is fixed to be softmax.

The loss functions supported are mean squared error and cross-entropy.

The implementation supports the following forms of gradient descent :

- Mini-batch gradient descent : Computes the gradient of m instances before making a parameter update
- Stochastic : Updates parameters with gradient for each instance. m has to be set to 1.
- Batch : Computes the gradient of the whole batch and then update gradient. m has to be set to the training set size.

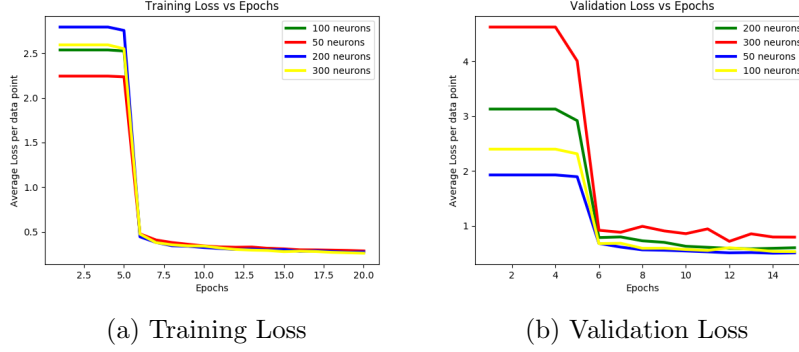


Figure 1: Single Layer Networks

The optimization algorithms supported are :

- Gradient Descent
- Momentum based gradient descent
- Nesterov based gradient descent
- Adam optimization

Annealing : When annealing of learning rate is enabled, the learning rate is reduced by half if the validation loss increases after the epoch.

In addition, we also implemented L2 regularization and tried dataset augmentation.

The dataset is first standardized, then trained on using an appropriate setting of the above hyperparameters, the learning rate is tuned accordingly, and the accuracy and F-score are scored on the validation set.

3 Sample Plots

3.1 Single Layer networks

Loss function : Cross-Entropy

Activation : Sigmoid

Optimization algorithm : Adam

Initial learning rate : 0.001

Batch size : 20

Figure 1 shows training and validation loss plots for 50, 100, 200 and 300 neurons.

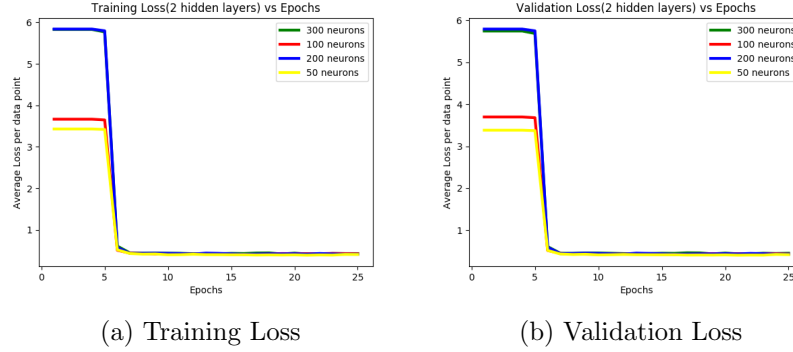


Figure 2: Two Layer Networks

3.2 Two Layer Networks

Loss function : Cross-Entropy
 Activation : Sigmoid
 Optimization algorithm : Adam
 Initial learning rate : 0.001
 Batch size : 20

Figure 2 shows training and validation loss plots for 50, 100, 200 and 300 neurons.

3.3 Three Layer Networks

Loss function : Cross-Entropy
 Activation : Sigmoid
 Optimization algorithm : Adam
 Initial learning rate : 0.001
 Batch size : 20

Figure 3 shows training and validation loss plots for 50, 100, 200 and 300 neurons.

3.4 Four Layer Network

Loss function : Cross-Entropy
 Activation : Sigmoid
 Optimization algorithm : Adam
 Initial learning rate : 0.001
 Batch size : 20

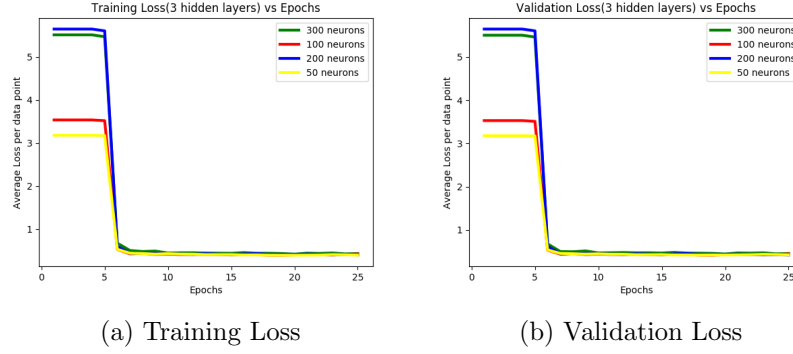


Figure 3: Three Layer Networks

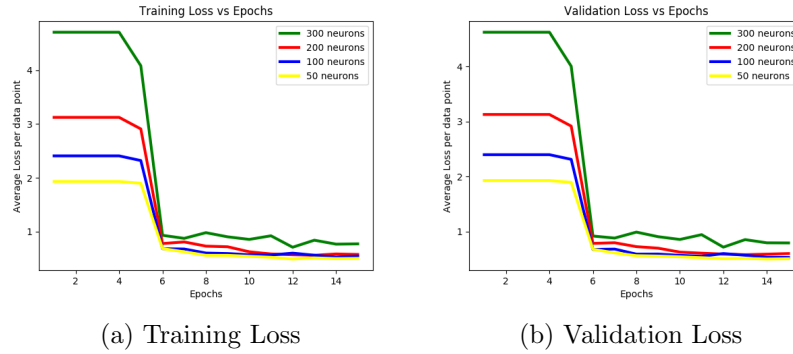


Figure 4: Four Layer Networks

Figure 4 shows training and validation loss plots for 50, 100, 200 and 300 neurons.

3.5 Adam, NAG, Momentum, GD

Network structure : 300, 300, 300

Loss function : Cross-Entropy

Activation : Sigmoid

Initial learning rate :

Batch size : 20

Figure 5 shows training and validation loss plots.




Figure 5: Adam, NAG, Momentum, GD

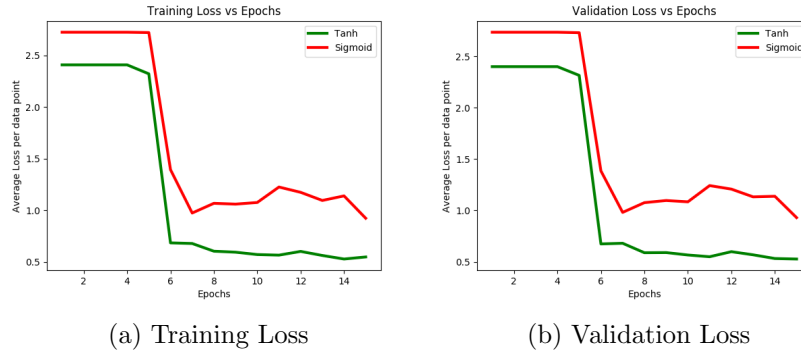


Figure 6: Sigmoid v/s Tanh activation


3.6 Sigmoid v/s Tanh activation

Network structure : 100, 100
 Loss function : Cross-Entropy
 Optimization algorithm : Adam
 Initial learning rate : 0.001
 Batch size : 20

Figure 6 shows training and validation loss plots.

3.7 Cross entropy loss v/s Squared error loss

Network structure : 100, 100
 Activation : Sigmoid
 Optimization algorithm : Adam



(a) Training Loss

(b) Validation Loss

Figure 7: Cross entropy loss v/s Squared error loss

Initial learning rate : 0.001

Batch size : 20

Figure 7 shows training and validation loss plots.

4 Regularisation Methods

We tried the following Regularisation Methods :

1. Bagging

Classifiers with best performance (F1- score) were bagged to give better performance. Models with best performance were trained and majority vote was given as class output. This helped achieve some improvement in performance.

2. Dataset Augmentation

- Data augmentation code for images was used from: <https://github.com/xkumiyu/numpy-data-augmentation>
- Vertical flip, Horizontal flip and random rotation were tried.
- Didn't result in a significant increase in performance.

.....

5 Results

We got the best results (on the validation as well as test data) with the following parameters b ensembling five different neural networks.

The five networks were :

- 100,100,100,100,100
- 100, 100, 50, 50, 50
- 200, 200, 200, 100, 100, 50
- 200,200,200,100,100,100
- 300,300,200,100,50

The parameters for each of these networks were :

Epochs : 20

Activation : Tanh

Optimization algorithm : NAG

Initial learning rate : 0.1

Batch size : 20

Loss function : Cross Entropy

The majority vote of these 5 networks was taken as the final prediction.